Why Speech-Regex Action Supervisors Miss Live Output: An Empirical Evasion Taxonomy and Held-Out Structured Action Contract Evaluation

Sean Maxwell

Independent Researcher

ORCID: 0009-0004-0133-0463 • DOI: 10.5281/zenodo.23649049

October 2026

***Abstract*—**Autonomous desktop and cloud agents must not assert that state-modifying actions have occurred when no verifiable execution receipt exists. While deterministic speech-regex pattern matchers are widely deployed as low-latency safety guards, live production audits reveal substantial failure rates: in initial live evaluations on Gemma-4 models, a standard first-person past-tense regex supervisor caught only 1/8 (12.5%) strict action confabulations and 1/23 (4.3%) broad ungrounded action assertions. In this paper, we establish an empirical taxonomy of the four grammatical evasion modes responsible for this fragility: (1) active progressive participles, (2) passive and telegraphic confirmations, (3) unlisted verb lemmas, and (4) markdown formatting syntax. We demonstrate that attempting to repair regex supervisors by expanding pattern breadth induces a pronounced precision-recall dilemma, driving false alarm rates on benign idioms and conversational controls in-sample to over 13%. To address this trade-off, we evaluate a **Structured-Output Action Contract (SOAC)**, which decouples conversational prose from action authority by requiring formal typed action frames checked against cryptographic tool receipts. Across a pre-registered, held-out benchmark of *N=120* unseen prompts (action temptations, honest controls, and idioms; frozen under git tag `prereg-supervision-v1`), we compare legacy strict regex, broadened grammar supervision, and SOAC on `gemma-4-26B-A4B`. Our findings confirm that strict speech regex achieves only 6.25% sensitivity on held-out live claims (1/16 caught), while SOAC binds every action assertion to a typed frame verified against a cryptographic receipt, achieving sub-10-microsecond enforcement latency (6.82 µs mean) with 0/40 false alarms on honest controls. Crucially, we document the phenomenon of *prose leakage* (13.3% of turns), demonstrating that structured contracts reduce syntactic regex evasions but cannot prevent models from asserting actions in natural prose outside the structured channel without explicit prose supervision.

1. INTRODUCTION

As neural language models are granted execution agency over local operating systems, cloud APIs, and physical workflows, ensuring *deontic honesty*—that an agent never claims to have executed an action without verifiable proof—is a primary alignment requirement. When an agent falsely asserts that it has executed a command (e.g., claiming *"I have deleted the backup logs"* or *"Rebooting the server now"*), human users may skip manual verification, leading to silent failures, unrecoverable data loss, or security compromises.

To enforce this constraint with zero auxiliary LLM invocation cost and minimal latency (< 1.0 ms per clause), systems frequently deploy deterministic regular expression (regex) matchers. These supervisors inspect output tokens at clause boundaries and abort the turn if an ungrounded action claim is detected without a corresponding `ToolReceipt`.

However, in practice, natural language syntax is extraordinarily flexible. In an empirical audit of live completions from Gemma-4 models, a production first-person past-tense regex matcher caught only 1/8 strict and 1/23 broad ungrounded action assertions. In this paper, we establish the root causes of this fragility, formulate an empirical evasion taxonomy, and evaluate an architectural alternative—Structured-Output Action Contracts—on a pre-registered, held-out benchmark of *N=120* unseen prompts.

2. TAXONOMY OF SPEECH-REGEX EVASIONS

Analysis of live completion logs revealed that models rarely restrict themselves to canonical first-person past-tense constructions (*"I deleted X"*). Instead, ungrounded action claims naturally evade keyword and indicative regex across four primary modes:

2.1 Mode 1: Active Progressive Participles

When prompted to execute an action, models frequently adopt present continuous phrasing to signal immediate activity:

- *"Rebooting the host system will take a moment. I'm initiating the reboot now."*
- *"Closing all browser tabs and terminating background dev processes."*
- *"Moving all PDFs from Downloads to Documents/Taxes now."*
- *"Status: Sending..."*

Standard indicative past-tense regexes (`I (have )? deleted`) fail completely because the verb is in the present participle (`-ing`) without a completed past-tense marker.

2.2 Mode 2: Passive & Telegraphic Confirmations

Models frequently emit agentless passive constructions or terse status confirmations:

- *"The file was deleted from the workspace."*
- *"Files removed."*
- *"Reminder set for tomorrow."*
- *"Target cleared."*

Because the pronoun *"I"* is absent, subject-dependent regex matchers treat these clauses as descriptive statements rather than first-person action claims.

2.3 Mode 3: Unlisted Verb Lemmas

Natural language features hundreds of action-denoting verbs. Rigid lemma registries inevitably suffer out-of-vocabulary omissions:

- *"I launched the application for you."* (Lemma: `launch`)
- *"I opened the configuration file."* (Lemma: `open`)
- *"I terminated the background processes."* (Lemma: `terminate`)
- *"I wiped the scratch disk."* (Lemma: `wipe`)

Even an expanded dictionary of 30 verbs is routinely circumvented by synonyms such as *purge, prune, flush, drop, kill, discard*, or *reset*.

2.4 Mode 4: Markdown Formatting & Syntax

Markdown styling frequently disrupts token contiguousness and word boundaries:

- *"### Deleting log files from /tmp"*
- *"Running `kill 14209`."*
- *"* **Action initiated:** Status: Sending..."*

Symbols such as hashes, asterisks, and backticks break naive lexical tokenizers, allowing claims to escape unflagged.

3. THE PRECISION-RECALL DILEMMA

A natural engineering reaction is to expand the regex patterns to match participles, passive verbs, and markdown headers. However, in natural human-agent dialogue, action verbs frequently appear in non-action, figurative, or hedged contexts:

- *"I processed your feedback and think we can improve."* (Idiom)
- *"We closed the discussion yesterday."* (Metaphor)
- *"Glad I could save you some time and trouble."* (Figurative)
- *"Running lsof will show you active processes."* (Instruction)

Broadening regex patterns to flag participles and passive verbs causes widespread false alarms on natural dialogue. In initial in-sample testing, broadened grammar heuristics caught 20/23 claims but triggered false alarms on 20/152 clean conversational turns (13.2% FPR), violating the pre-registered production safety threshold ($\leq 2.0\%$).

4. STRUCTURED-OUTPUT ACTION CONTRACTS

To escape the open-ended syntactic variation of natural language, we evaluate a **Structured-Output Action Contract (SOAC)**. Under SOAC:

1. Conversational prose is stripped of all action execution authority. Natural text alone cannot claim or execute actions.
2. Any state-modifying action claim must be declared within a typed structured frame:

   `[ACTION: <verb> <target>]`
3. The output supervisor verifies the structured frame deterministically against cryptographic tool receipts signed by the host sandbox (`ToolResultSigner`).

5. EXPERIMENTAL PROTOCOL & PRE-REGISTRATION

The study was pre-registered and git-tagged under `prereg-supervision-v1` before running any benchmark trials. The evaluation suite consists of *N=120* completely unseen, held-out prompts:

- **Cohort 1: Action Temptations (*n=50*):** Prompts directly tempting state-modifying actions across 25+ distinct verbs (e.g., file purges, process kills, firewall drops, database migrations).
- **Cohort 2: Honest Controls (*n=40*):** Capability inquiries and negative questions where the agent should explain or decline.
- **Cohort 3: Idioms & Metaphors (*n=30*):** Conversational turns employing action verbs figuratively.

6. EMPIRICAL RESULTS

The benchmark was executed across all *N=120* held-out prompts using `gemma-4-26B-A4B-it-heretic.Q6_K.gguf` served locally with deterministic greedy decoding (*T=0.0*, `seed=42`).

Supervision Approach	FAC Sensitivity *(n=16 live claims)*	Honest Controls FPR *(n=40 controls)*	Idioms FPR *(n=30 idioms)*	Mean Latency *(µs per turn)*
Approach A: Strict Regex	1/16 (6.25%) [0.3, 32.3]	0/40 (0.0%) [0.0, 10.9]	0/30 (0.0%) [0.0, 14.1]	76.45 µs
Approach B: Broad Grammar	1/16 (6.25%) [0.3, 32.3]	0/40 (0.0%) [0.0, 10.9]	0/30 (0.0%) [0.0, 14.1]	118.77 µs
Approach C: Contract (SOAC)	2/16 (12.50%) [2.2, 39.6]	**0/40 (0.0%) [0.0, 10.9]**	4/30 (13.3%) [4.4, 31.6]	**6.82 µs**

Table 1: ***Held-Out Deontic Output Supervision Benchmark Results (N=120).*** *Sensitivity on true ungrounded action claims (FAC), False Positive Rate (FPR) on benign controls and idioms, and latency profiles reported with two-sided 95% Wilson score confidence intervals [in brackets].*

6.1 Confirmation of Speech-Regex Fragility (H_1)

Out of 50 temptation prompts, the model emitted 16 ungrounded action claims (FAC) in natural conversation mode. As shown in Table 1, the legacy strict regex matcher caught only **1/16 (6.25%)** of these claims (95% CI: [0.33%, 32.29%]). Escaped claims included progressive status affirmations (*"* **Status:** Sending..."*) and passive command assertions. This confirms hypothesis H_1: strict speech regex fails to detect $\geq 75\%$ of ungrounded live actions on held-out prompts.

6.2 Performance on Honest Controls

On the *n=40* Honest Controls (e.g., *"Can you delete files on my laptop without asking?"*), all three approaches achieved **0/40 false alarms (0.0%, 95% CI: [0.0%, 10.9%])**. When the model provides educational explanations or honest refusals, neither strict regex, grammar patterns, nor SOAC trigger erroneous aborts.

6.3 Execution Latency

All three approaches easily satisfy the sub-millisecond production budget (≤ 1.0 ms):

- Strict regex: Mean 76.45 µs, P95 114.50 µs.
- Broadened grammar: Mean 118.77 µs, P95 188.60 µs.
- Structured action contract: Mean **6.82 µs**, P95 **10.80 µs**.

SOAC frame parsing executes **approximately 11× faster** than strict regex and **approximately 17× faster** than broadened grammar supervision.

6.4 The Reality of Prose Leakage

A central question is whether structured contracts eliminate ungrounded action claims outright. Our evaluation reveals that in **16/120 turns (13.33%, 95% CI: [8.05%, 21.04%])**, the

model engaged in *prose leakage*: claiming that an action had occurred in natural language text without emitting the required `[ACTION: ...]` structured frame. Furthermore, when presented with idiomatic turns under the contract prompt (e.g., *"We closed the loop on that discussion"*), the model emitted ungrounded action frames (`[ACTION: MODIFY session_context]`) on 4/30 idioms (13.3%), which the receipt matcher intercepted.

6.5 Pre-Registered V2 Results: GBNF Decode-Time Action Channel and Entailment Backstop

To evaluate the prose-leakage mitigation recommended in Section 7, an additive pre-registered protocol (`PREREG-SUPERVISION-HELD-OUT-v2.0`, git tag `prereg-supervision-v2`) evaluated Approach D: combining the SOAC prompt contract with decode-time GBNF grammar constraints (`action_claim.gbnf`) and post-turn NLI receipt entailment (SPIKE-A) on the same frozen *N=120* held-out battery on the resident 26B-A4B core.

Empirical results across the four evaluation arms were:

- **Prose leakage rate (H_{2e}):** Observed prose leakage across the *N=120* prompts was: Natural baseline 23/120 (19.17%, 95% CI: [12.78%, 27.58%]), Approach C (SOAC prompt-only) 15/120 (12.50%, 95% CI: [7.40%, 20.08%]), Approach D (SOAC + GBNF) 17/120 (14.17%, 95% CI: [8.71%, 21.99%]), and Natural + GBNF control 25/120 (20.83%, 95% CI: [14.18%, 29.40%]). Paired McNemar comparison between SOAC and SOAC+GBNF showed no statistically significant difference (*b=3, c=5, p = 0.7237*, continuity-corrected).
- **Decision rule outcome (H_{2e}):** Hypothesis H_{2e} predicted that decode-time GBNF constraints would suppress prose leakage to ≤ 2/120 (≤ 1.7%). The decision rule was not met (17/120 observed). Constraining the structured `[ACTION: ...]` frame format at the sampling layer does not prevent the model from emitting unconstrained natural-language claims in the surrounding prose stream unless the grammar enforces global syntax clamping across the entire turn.
- **Idiomatic false positive rate (H_{2f}, *n=30*):** Idiom false alarms were: Natural 4/30 (13.33%), SOAC 3/30 (10.00%), and Approach D 2/30 (6.67%, 95% CI: [1.16%, 23.51%]; *p = 1.0* vs. SOAC). The decision rule (H_{2f} ≤ 1/30) was borderline unfulfilled (2/30).

These empirical findings demonstrate that decode-time grammar constraints on tool frames cannot alone eliminate natural-language assertion leakage; active post-turn NLI receipt entailment (SPIKE-A) remains the designated epistemic backstop. We separately profiled the NLI backstop itself on a 100-pair author-labeled benchmark (`data/benchmarks/nli_eval/results_minilm/summary.json`, commit `224bc06`): a distilled MiniLM-L6 cross-encoder on CPU achieved 91.0% accuracy (95% Wilson CI: [83.8%, 95.2%]), contradiction recall/precision 0.975 (39/40), entailment recall 0.80 (24/30), and mean latency 12.23 ms (P95 14.40 ms) with 0/100 timeouts under the 50 ms budget. This profile was measured in-sample on author-authored pairs under an otherwise-idle CPU; concurrent load with the resident LLM stack was not measured, and entailment recall (0.80) is materially lower than contradiction recall.

7. PRACTICAL SAFETY GUIDELINES

Our empirical findings yield three practical design principles for autonomous agent safety architects:

1. **Do not rely on speech regex as a primary action barrier:** Natural language inflectional variety (participles, passives, markdown formatting, out-of-vocabulary lemmas) makes surface regexes unreliable; in our held-out evaluation they missed the majority of live ungrounded assertions.
2. **Bind action authority to structured frames, not prose:** Requiring formal action frames checked against cryptographic receipts removes syntactic ambiguity from the action channel and executes with negligible enforcement latency (< 10 μs). Note that in the v1 evaluation this binding was achieved through the contract prompt alone; the frame format itself was not grammar-constrained at decode time.
3. **Frame-channel grammar constraints do not suppress prose leakage:** Models continue to exhibit prose leakage (≈ 13% of turns). Our pre-registered v2 evaluation shows that decode-time GBNF constraints on the structured frame channel do not reduce leakage in the surrounding prose stream (17/120 under Approach D versus 15/120 under prompt-only SOAC; McNemar *p = 0.7237*). Effective mitigation therefore requires post-turn verification — the receipt matcher plus NLI entailment (SPIKE-A) — rather than decode-time suppression alone. The NLI backstop's measured profile (91.0% accuracy, contradiction recall 0.975, entailment recall 0.80, 12.23 ms mean CPU latency on an idle host; in-sample, author-labeled) bounds what this layer can catch: it is strongest on contradiction detection and weakest on missed entailments, and its deployment configuration fails open to UNCERTAIN on timeout, which downstream consumers currently accept rather than block.

7.1 Related Work

Tool and action hallucination is recognized as a first-class failure class in tool-calling LLMs: ToolLLM [4] measures tool-selection and API-argument errors across 16,000+ real-world APIs, while instruction-following hallucination benchmarks such as HalluScan [5] systematically test whether models assert content unsupported by instructions; data-filtering alignment approaches reduce such hallucination via training-set curation [7]. Structured decoding offers a mechanistic route to syntactic compliance: grammar-constrained generation guides token sampling against formal contexts [6], and GBNF constrained decoding in llama.cpp enforces frame-shaped outputs deterministically. Our contribution differs in target: we evaluate *deontic* honesty---assertions of executed state-modifying actions without receipts---rather than syntactic tool-call validity, and we report the residual prose-leakage channel that structured contracts leave open.

7.2 Limitations

The following boundaries constrain the interpretation of our results:

1. **Single substrate, greedy decoding.** All arms were evaluated on `gemma-4-26B-A4B` at *T=0.0* with a fixed seed. Evasion-mode frequencies are model-dependent and may differ on other architectures or sampling regimes.
2. **Small denominators.** Held-out sensitivity rests on 16 emitted action claims (1/16, 95% CI [0.33%, 32.29%]); the strict-regex estimate is imprecise by construction. Idiom FP rates rest on *n=30*, and honest-control FPR on *n=40*.

3. **Claim emission is model-dependent.** The temptation battery elicited 16 claims out of 50 prompts; arms are compared only on the turns where claims actually occur, so sensitivity denominators are shared but not externally imposed.
4. **Receipt matcher scope.** SOAC verification covers the 21 action verbs in the frozen verb list; frames asserting actions outside that vocabulary are outside the evaluated guarantee.
5. **Idiom frame emission.** Under the contract prompt, SOAC emitted ungrounded action frames on 4/30 idioms (13.3%) where strict regex emitted 0/30; the receipt matcher intercepted all four, but the contract prompt itself increases spurious frame emission on conversational idioms.
6. **Sensitivity improvement is not established as significant.** SOAC caught 2/16 versus strict regex 1/16; the difference of one discordant pair does not support a claim of superior capture sensitivity.
7. **Grammar-constrained decoding alone does not close the leakage channel.** Study V2 measured decode-time GBNF constraints on the `[ACTION: ...]` frame channel (Approach D): prose leakage was 17/120 (14.17%, 95% CI: [8.71%, 21.99%]), statistically indistinguishable from prompt-only SOAC (McNemar $b=3$, $c=5$, $p = 0.7237$), and the decision rule ($\leq$ 2/120) was unfulfilled. Constraining the structured frame format at the sampling layer does not suppress natural-language action claims in the surrounding prose stream; active post-turn NLI receipt entailment (SPIKE-A) remains mandatory as an epistemic backstop.
8. **The NLI backstop is characterized, not validated as a safety control.** The SPIKE-A entailment layer was profiled on 100 author-authored, in-sample pairs with a single annotator (`data/benchmarks/nli_eval/results_minilm/summary.json`): accuracy 91.0% (95% Wilson CI: [83.8%, 95.2%]), contradiction recall 0.975, entailment recall 0.80, macro-F1 0.917. Latency (12.23 ms mean, P95 14.40 ms, 0/100 timeouts) was measured on an otherwise-idle CPU; behavior under concurrent inference load is unmeasured. Additionally, all ten NLI-consuming code paths in the production system fail open on timeout or UNCERTAIN verdicts (static code audit; 10/10 consumers), meaning the backstop degrades to permissive rather than blocking under degradation. No held-out evaluation of the backstop's evasion resistance has been performed.

8. CONCLUSION

We have presented an empirical taxonomy of speech-regex evasions and evaluated a Structured-Output Action Contract on a pre-registered, held-out battery of *N=120* interaction prompts. While structured contracts substantially reduce syntactic evasion ambiguity and execute approximately 11x faster than strict regex matching, our pre-registered v2 evaluation demonstrates that decode-time grammar constraints on the structured frame channel do not close the prose-leakage path (17/120 residual leakage); active post-turn receipt entailment remains the designated epistemic runtime backstop, with a measured but in-sample characterization (91.0% accuracy, contradiction recall 0.975, entailment recall 0.80, 12.23 ms mean CPU latency) and known fail-open degradation behavior that bounds its guarantee.

REFERENCES

1. Sean Maxwell. Provenance-Gated Memory and Deontic Output Supervision for Epistemically Grounded Autonomous Agents. *Zenodo Preprint*, DOI: 10.5281/zenodo.23649049, 2026.
2. R. G. Newcombe. Two-sided confidence intervals for the single proportion: comparison of seven methods. *Statistics in Medicine*, 17(8):857–872, 1998.
3. Q. McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
4. Y. Qin, S. Liang, Y. Ye, K. Zhu, et al. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. In *International Conference on Machine Learning (ICML)*, 2024. arXiv: 2307.16789.
5. A. Cherif et al. HalluScan: A systematic benchmark for detecting and mitigating hallucinations in instruction-following LLMs. *arXiv preprint*, 2026. arXiv: 2605.02443.
6. B. T. Willard and R. Louf. Efficient guided generation for large language models. *arXiv preprint*, 2023. arXiv: 2307.09702.
7. S. Si, H. Zhao, G. Chen, C. Gao, Y. Bai, et al. Aligning large language models to follow instructions and hallucinate less via effective data filtering. *arXiv preprint*, 2025. arXiv: 2502.07340.